

# Range Minimum Query: A Friendly Step-by-Step Tutorial

A Comprehensive Guide to RMQ Algorithms

2025

## Contents

What is Range Minimum Query (RMQ)?	3
Algorithm 1: Naive Approach (The Straightforward Way)	4
How It Works	4
Step-by-Step Example	4
Visual Representation	4
Time Complexity Deep Dive	4
Algorithm 2: Dynamic Programming (Pre-calculate Everything!)	6
The Idea	6
Building the DP Table Step-by-Step	6
Final DP Table	6
Query Example	6
Time Complexity Deep Dive	6
Algorithm 3: Sparse Table (Binary Lifting Magic!)	8
The Clever Idea	8
Building the Sparse Table Step-by-Step	8
Sparse Table Visualization	8
Query Example: Query(2, 6)	9
Visual of Query Coverage	9
Time Complexity Deep Dive	9
Algorithm 4: Block Decomposition (Square Root Decomposition)	12
The Balanced Idea	12
Building Blocks Step-by-Step	12
Visual Representation	12
Query Example: Query(1, 7)	12
Visual of Query	12
Time Complexity Deep Dive	13
Algorithm 5: LCA-based RMQ (The Tree Transformation!)	15
The Amazing Connection	15
Building the Cartesian Tree Step-by-Step	15
Final Cartesian Tree	16
How RMQ becomes LCA	16
LCA using Binary Lifting	17
Time Complexity Deep Dive	17
Comparison: Which Algorithm to Choose?	20

Quick Decision Guide . . . . .	20
Performance Summary Table . . . . .	20
Interactive Examples . . . . .	21
Let's Trace Through a Complete Example . . . . .	21
Practice Problems . . . . .	22
Problem 1: Build Your Own DP Table . . . . .	22
Problem 2: Sparse Table Query . . . . .	22
Tips and Tricks . . . . .	23
1. Sparse Table Power-of-2 Trick . . . . .	23
2. Block Size Selection . . . . .	23
3. DP Memory Optimization . . . . .	23
4. Cartesian Tree Stack Trick . . . . .	23
Deep Dive: Understanding Complexity Growth . . . . .	24
How Complexities Compare As N Grows . . . . .	24
Time Complexity Visualization . . . . .	24
When Each Algorithm Wins . . . . .	24
Memory vs Speed Trade-offs . . . . .	25
Big-O Doesn't Tell the Whole Story! . . . . .	25
Amortized Analysis: When Average Case Matters . . . . .	25
The Complexity Hierarchy . . . . .	26
Conclusion . . . . .	26

## What is Range Minimum Query (RMQ)?

Imagine you have a list of numbers, and someone keeps asking you: “What’s the smallest number between position 3 and position 7?” That’s exactly what RMQ solves!

### Example Array:

Index:	0	1	2	3	4	5	6	7
Value:	5	2	4	7	1	3	6	8

**Query(2, 5)** = “What’s the minimum between index 2 and 5?” - Look at values: [4, 7, 1, 3] - Answer: 1 (at index 4)

Now, let’s explore 5 different ways to solve this problem, from simple to sophisticated!

## Algorithm 1: Naive Approach (The Straightforward Way)

### How It Works

Just look at every element in the range and find the minimum. It's like reading through a list with your finger!

### Step-by-Step Example

**Array:** [5, 2, 4, 7, 1, 3, 6, 8]

**Query(2, 5):**

Step 1: Look at index 2 → value is 4, min = 4

Step 2: Look at index 3 → value is 7, min = 4 (no change)

Step 3: Look at index 4 → value is 1, min = 1 (new minimum!)

Step 4: Look at index 5 → value is 3, min = 1 (no change)

Answer: 1

### Visual Representation

Query(2, 5):

[5, 2, |4, 7, 1, 3|, 6, 8]

↑ ↑ ↑ ↑

Check each one

Return: 1

### Time Complexity Deep Dive

#### Why is Preprocessing $O(1)$ ?

- We literally do nothing! Just store the array as-is.
- No computation, no extra data structures.
- Time taken: constant, regardless of array size.

**Why is Query Time  $O(n)$ ?** Let's count the operations:

Query(L, R):

1. Initialize min = array[L] → 1 operation

2. For each element from L+1 to R:

- Compare with current min → (R-L) comparisons

- Update min if needed → up to (R-L) assignments

Total:  $2(R-L) + 1$  operations

**Worst case:** Query(0, n-1) checks all n elements →  $O(n)$  **Best case:** Query(i, i) checks 1 element →  $O(1)$  **Average case:** Query covers  $n/2$  elements →  $O(n)$

### Space Complexity: $O(n)$

- Original array: n elements × 4 bytes (for int) = 4n bytes
- No additional structures needed
- Total space:  $O(n)$

**Real-world Performance** For an array of 1,000,000 elements: - Preprocessing: 0 microseconds  
- Query (worst case): ~1,000 microseconds (1 ms) - Query (average): ~500 microseconds

**When to Use:** - Queries are rare ( $< 100$  queries total) - Array changes frequently (after every few queries) - Array is small ( $< 1000$  elements)

## Algorithm 2: Dynamic Programming (Pre-calculate Everything!)

### The Idea

What if we pre-calculate the answer for EVERY possible range? Then queries become instant lookups!

### Building the DP Table Step-by-Step

**Array:** [5, 2, 4, 7]

We'll build a table where  $dp[i][j]$  = minimum value from index  $i$  to  $j$ .

#### Step 1: Single elements (length = 1)

```
dp[0][0] = 5 (just element at index 0)
dp[1][1] = 2 (just element at index 1)
dp[2][2] = 4 (just element at index 2)
dp[3][3] = 7 (just element at index 3)
```

#### Step 2: Pairs (length = 2)

```
dp[0][1] = min(5, 2) = 2
dp[1][2] = min(2, 4) = 2
dp[2][3] = min(4, 7) = 4
```

#### Step 3: Triples (length = 3)

```
dp[0][2] = min(dp[0][1], 4) = min(2, 4) = 2
dp[1][3] = min(dp[1][2], 7) = min(2, 7) = 2
```

#### Step 4: Full array (length = 4)

```
dp[0][3] = min(dp[0][2], 7) = min(2, 7) = 2
```

### Final DP Table

	j→	0	1	2	3
i↓		-----			
0		5	2	2	2
1		-	2	2	2
2		-	-	4	4
3		-	-	-	7

### Query Example

**Query(1, 3):** Just look up  $dp[1][3] = 2$ . Instant!

### Time Complexity Deep Dive

**Why is Preprocessing  $O(n^2)$ ?** Let's count exactly how many cells we fill:

For array of size  $n$ :

- Ranges of length 1:  $n$  cells
- Ranges of length 2:  $n-1$  cells
- Ranges of length 3:  $n-2$  cells
- ...
- Ranges of length  $n$ : 1 cell

Total cells =  $n + (n-1) + (n-2) + \dots + 1 = n(n+1)/2$

**Mathematical proof:**

Sum =  $n(n+1)/2 = (n^2 + n)/2 = O(n^2)$

**Actual operations per cell:**

$dp[i][j] = \min(dp[i][j-1], array[j])$

↑ 1 lookup + 1 comparison + 1 assignment = 3 operations

Total operations =  $3 \times n^2/2 = O(n^2)$

**Why is Query Time  $O(1)$ ?**

Query(L, R):

1. Access  $dp[L][R]$  → 1 array access

2. Return value → 1 operation

Total: 2 operations =  $O(1)$

No loops, no comparisons, just direct memory access!

**Space Complexity:  $O(n^2)$**

For array of size  $n$ :

- DP table:  $n \times n \times 4$  bytes =  $4n^2$  bytes
  - Original array:  $n \times 4$  bytes =  $4n$  bytes
- Total:  $4n^2 + 4n = O(n^2)$

**Memory usage examples:** -  $n = 100$ : ~40 KB -  $n = 1,000$ : ~4 MB -  $n = 10,000$ : ~400 MB (getting expensive!) -  $n = 100,000$ : ~40 GB (impractical!)

**Building Time Analysis** For an array of size  $n$ :

for (length = 1 to  $n$ ): →  $n$  iterations  
    for (start = 0 to  $n$ -length): → average  $n/2$  iterations  
         $dp[start][end] = \dots$  →  $O(1)$  operation  
Total:  $n \times n/2 \times 1 = O(n^2)$

**Real-world Performance** For  $n = 1,000$ : - Preprocessing: ~2-3 milliseconds - Memory used: ~4 MB - Query time: ~0.01 microseconds (10 nanoseconds!) - Break-even point: Need ~2,000 queries to justify preprocessing

**When to Use:** - Array size  $< 2,000$  elements - Number of queries  $> n^2/1000$  - Can afford  $O(n^2)$  memory - Static data (no updates)

## Algorithm 3: Sparse Table (Binary Lifting Magic!)

### The Clever Idea

Instead of storing ALL ranges, only store ranges with lengths that are powers of 2 (1, 2, 4, 8, ...). Any range can be covered by at most 2 overlapping power-of-2 ranges!

### Building the Sparse Table Step-by-Step

**Array:** [5, 2, 4, 7, 1, 3, 6, 8]

We build  $st[i][j]$  = minimum in range starting at  $i$  with length  $2^j$ .

#### Step 1: Length 1 ( $2^0 = 1$ )

```
st[0][0] = 5 (range [0,0])
st[1][0] = 2 (range [1,1])
st[2][0] = 4 (range [2,2])
st[3][0] = 7 (range [3,3])
st[4][0] = 1 (range [4,4])
st[5][0] = 3 (range [5,5])
st[6][0] = 6 (range [6,6])
st[7][0] = 8 (range [7,7])
```

#### Step 2: Length 2 ( $2^1 = 2$ )

```
st[0][1] = min(st[0][0], st[1][0]) = min(5, 2) = 2 (range [0,1])
st[1][1] = min(st[1][0], st[2][0]) = min(2, 4) = 2 (range [1,2])
st[2][1] = min(st[2][0], st[3][0]) = min(4, 7) = 4 (range [2,3])
st[3][1] = min(st[3][0], st[4][0]) = min(7, 1) = 1 (range [3,4])
st[4][1] = min(st[4][0], st[5][0]) = min(1, 3) = 1 (range [4,5])
st[5][1] = min(st[5][0], st[6][0]) = min(3, 6) = 3 (range [5,6])
st[6][1] = min(st[6][0], st[7][0]) = min(6, 8) = 6 (range [6,7])
```

#### Step 3: Length 4 ( $2^2 = 4$ )

```
st[0][2] = min(st[0][1], st[2][1]) = min(2, 4) = 2 (range [0,3])
st[1][2] = min(st[1][1], st[3][1]) = min(2, 1) = 1 (range [1,4])
st[2][2] = min(st[2][1], st[4][1]) = min(4, 1) = 1 (range [2,5])
st[3][2] = min(st[3][1], st[5][1]) = min(1, 3) = 1 (range [3,6])
st[4][2] = min(st[4][1], st[6][1]) = min(1, 6) = 1 (range [4,7])
```

#### Step 4: Length 8 ( $2^3 = 8$ )

```
st[0][3] = min(st[0][2], st[4][2]) = min(2, 1) = 1 (range [0,7])
```

### Sparse Table Visualization

	Length→			
Index↓	1	2	4	8
0	5	2	2	1



1	2	2	1	-
2	4	4	1	-
3	7	1	1	-
4	1	1	1	-
5	3	3	-	-
6	6	6	-	-
7	8	-	-	-

### Query Example: Query(2, 6)

Range length =  $6 - 2 + 1 = 5$

Largest power of 2  $\leq 5$  is 4 ( $2^2$ )

Split into two overlapping ranges of length 4:

- Range 1: [2, 5]  $\rightarrow$  st[2][2] = 1

- Range 2: [3, 6]  $\rightarrow$  st[3][2] = 1

Answer:  $\min(1, 1) = 1$

### Visual of Query Coverage

Query [2, 6]:

Index: 0 1 2 3 4 5 6 7

Value: 5 2 4 7 1 3 6 8

-----	(st[2][2]: covers 2-5)
-----	(st[3][2]: covers 3-6)
=====	(Full coverage with overlap!)

### Time Complexity Deep Dive

Why is Preprocessing  $O(n \log n)$ ? Understanding  $\log n$  levels:

For array of size  $n = 16$ :

- Level 0 (length 1): 16 entries
- Level 1 (length 2): 15 entries
- Level 2 (length 4): 13 entries
- Level 3 (length 8): 9 entries
- Level 4 (length 16): 1 entry

Number of levels =  $\log(16) + 1 = 5$  levels

Counting total operations:

```

for j from 0 to log(n):            $\rightarrow$  log n iterations
  for i from 0 to  $n - 2^j$ :          $\rightarrow$   $(n - 2^j + 1)$  iterations
    st[i][j] = min(st[i][j-1], st[i+2^(j-1)][j-1])
                     $\uparrow$  2 lookups + 1 comparison + 1 assignment

```

Total entries =  $n \times 1 + (n-1) \times 1 + (n-3) \times 1 + \dots$

$n \times \log n$  entries

Each entry:  $O(1)$  operation  
Total:  $O(n \log n)$

### Mathematical analysis:

Sum =  $\sum_{j=0}^{\log n} (n - 2^j + 1)$   
=  $n \times \log(n) - (2^{(\log n + 1)} - 1) + \log(n)$   
=  $n \times \log(n) - (2n - 1) + \log(n)$   
=  $O(n \log n)$

### Why is Query Time $O(1)$ ? The brilliant trick:

Query(L, R):

1. Calculate  $k = \text{floor}(\log(R - L + 1)) \rightarrow O(1)$  with bit operations
  2. Access  $st[L][k] \rightarrow O(1)$  array access
  3. Access  $st[R - 2^k + 1][k] \rightarrow O(1)$  array access
  4. Return min of the two  $\rightarrow O(1)$  comparison
- Total: 4 operations =  $O(1)$

### How to calculate $\log$ in $O(1)$ :

```
// Using built-in functions (compiled to single CPU instruction)  
int k = __builtin_clz(1) - __builtin_clz(R - L + 1);  
// or  
int k = 31 - __builtin_clz(R - L + 1); // for 32-bit integers
```

### Space Complexity: $O(n \log n)$

Sparse table dimensions:

- Rows:  $n$  (one for each starting position)
- Columns:  $\log(n) + 1$  (one for each power of 2)
- Each cell: 4 bytes (integer)

Total space =  $n \times (\log n + 1) \times 4 \text{ bytes} = O(n \log n)$

**Memory usage examples:** -  $n = 1,000$ : ~40 KB ( $\log(1000) = 10$ ) -  $n = 100,000$ : ~6.4 MB ( $\log(100000) = 17$ ) -  $n = 1,000,000$ : ~80 MB ( $\log(1000000) = 20$ )

Much better than DP's  $O(n^2)$ !

### Why Can We Overlap Ranges? This only works because MIN is an **idempotent** operation:

$\min(a, a) = a$   
 $\min(\min(a, b), \min(b, c)) = \min(a, b, c)$

So overlapping doesn't affect the result!

**Real-world Performance** For  $n = 100,000$ : - Preprocessing: ~15-20 milliseconds - Memory used: ~6.4 MB - Query time: ~0.05 microseconds (50 nanoseconds) - Can handle millions of queries per second!

**Comparison with DP:** - DP for  $n=100,000$ : 40 GB memory (impractical) - Sparse Table: 6.4 MB memory (very practical) - Both have  $O(1)$  query, but Sparse Table scales much better

**When to Use:** - Large static arrays (up to  $10^6$  elements) - Need absolutely fastest query time - Can afford  $O(n \log n)$  preprocessing - No updates to the array

## Algorithm 4: Block Decomposition (Square Root Decomposition)

### The Balanced Idea

Divide the array into blocks of size  $\sqrt{n}$ . Pre-compute the minimum for each complete block. For queries, combine partial blocks with complete blocks.

### Building Blocks Step-by-Step

**Array:** [5, 2, 4, 7, 1, 3, 6, 8, 9] ( $n = 9$ ,  $\text{block\_size} = 3$ )

#### Step 1: Divide into blocks

Block 0: [5, 2, 4]

Block 1: [7, 1, 3]

Block 2: [6, 8, 9]

#### Step 2: Pre-compute block minimums

$\text{block\_min}[0] = \min(5, 2, 4) = 2$

$\text{block\_min}[1] = \min(7, 1, 3) = 1$

$\text{block\_min}[2] = \min(6, 8, 9) = 6$

### Visual Representation

Array:	[5, 2, 4]	[7, 1, 3]	[6, 8, 9]
	Block 0	Block 1	Block 2
Min:	2	1	6

### Query Example: Query(1, 7)

#### Step 1: Identify affected blocks

Index 1 is in Block 0 (partial)

Index 7 is in Block 2 (partial)

Block 1 is completely covered

#### Step 2: Calculate minimum

Partial Block 0: elements [1, 2]  $\rightarrow \min(2, 4) = 2$

Complete Block 1:  $\text{block\_min}[1] = 1$

Partial Block 2: element [6, 7]  $\rightarrow \min(6, 8) = 6$

Answer:  $\min(2, 1, 6) = 1$

### Visual of Query

Query [1, 7]:

[5,  2, 4]	[7, 1, 3]	[6, 8 , 9]
↑~~~~~	~~~~~	~~~~~↑
partial	complete	partial

## Time Complexity Deep Dive

### Why is Preprocessing $O(n)$ ?

Preprocessing steps:

1. Determine block size =  $\sqrt{n}$   $\rightarrow O(1)$
2. Create block\_min array of size  $\sqrt{n}$   $\rightarrow O(\sqrt{n})$
3. For each element in array:  $\rightarrow n$  iterations
  - Assign to a block  $\rightarrow O(1)$
  - Update block minimum  $\rightarrow O(1)$

Total:  $O(1) + O(\sqrt{n}) + n \times O(1) = O(n)$

Detailed breakdown:

```
block_size = sqrt(n);           // O(1)
for (int i = 0; i < n; i++) {   // n iterations
    int block_id = i / block_size; // O(1) - integer division
    block_min[block_id] = min(block_min[block_id], arr[i]); // O(1)
}
Total: O(n)
```

### Why is Query Time $O(\sqrt{n})$ ? Three parts of any query:

1. **Left partial block:** Up to  $\sqrt{n} - 1$  elements
2. **Complete middle blocks:** Up to  $\sqrt{n} - 2$  blocks (each takes  $O(1)$  to check)
3. **Right partial block:** Up to  $\sqrt{n} - 1$  elements

Worst case analysis:

- Left partial:  $\sqrt{n} - 1$  comparisons
- Middle blocks:  $\sqrt{n} - 2$  lookups
- Right partial:  $\sqrt{n} - 1$  comparisons

Total:  $(\sqrt{n} - 1) + (\sqrt{n} - 2) + (\sqrt{n} - 1) = 3\sqrt{n} - 4 = O(\sqrt{n})$

Example with  $n = 100$  (block\_size = 10):

Query(5, 84):

- Left partial [5-9]: 5 elements  $\rightarrow 5$  operations
- Middle blocks [10-79]: 7 blocks  $\rightarrow 7$  operations
- Right partial [80-84]: 5 elements  $\rightarrow 5$  operations

Total: 17 operations  $1.7\sqrt{n}$

### Why Block Size = $\sqrt{n}$ is Optimal? Let's analyze with block size = $b$ :

- Number of blocks:  $n/b$
- Elements per block:  $b$
- Query time:  $O(b)$  for partials +  $O(n/b)$  for complete blocks

Total query time:  $O(b + n/b)$

To minimize, take derivative and set to 0:

$$d/db (b + n/b) = 1 - n/b^2 = 0$$

$$b^2 = n$$

$$b = \sqrt{n}$$

**What if we use different block sizes?** - Block size =  $n/10$ : Query =  $O(n/10)$ , not good! - Block size = 10: Query =  $O(n/10)$ , still linear! - Block size =  $\sqrt{n}$ : Query =  $O(\sqrt{n})$ , perfectly balanced!

**Space Complexity:**  $O(\sqrt{n} + n)$

Storage requirements:

- Original array:  $n$  elements  $\times$  4 bytes =  $4n$  bytes
  - Block minimums:  $\sqrt{n}$  blocks  $\times$  4 bytes =  $4\sqrt{n}$  bytes
  - Block boundaries:  $2\sqrt{n}$  integers =  $8\sqrt{n}$  bytes (optional)
- Total:  $O(n + \sqrt{n}) = O(n)$

But we often say  $O(\sqrt{n})$  for *additional* space beyond the input array.

**Update Complexity:  $O(1)$**  This is where Block Decomposition shines!

Update(index, new\_value):

1. Update array[index] = new\_value  $\rightarrow O(1)$
  2. Find block\_id = index / block\_size  $\rightarrow O(1)$
  3. Recompute minimum for that block  $\rightarrow O(\sqrt{n})$
- OR
- If new\_value < block\_min: update  $\rightarrow O(1)$
- If old\_value was min: recompute  $\rightarrow O(\sqrt{n})$

Smart update strategy: - If new value is smaller than block min:  $O(1)$  - Otherwise:  $O(\sqrt{n})$  to recompute one block - Compare to other algorithms: - Sparse Table:  $O(n \log n)$  rebuild everything! - DP:  $O(n^2)$  rebuild everything!

**Real-world Performance** For  $n = 1,000,000$  (block\_size = 1,000): - Preprocessing: ~1 millisecond - Memory used: ~4 KB extra (just block minimums) - Query time: ~1 microsecond - Update time: ~1 microsecond (average case)

**Performance comparison:**

Algorithm	Query	Update	Memory
Naive	1000 s	0.01 s	0
Block Decomp	1 s	1 s	4 KB
Sparse Table	0.05 s	20,000 s	80 MB

**Mathematical Beauty: The Square Root Appears Everywhere!** For array of size  $n$  with block size  $\sqrt{n}$ : - Number of blocks:  $\sqrt{n}$  - Elements per block:  $\sqrt{n}$  - Query touches at most:  $2\sqrt{n}$  elements +  $\sqrt{n}$  blocks =  $3\sqrt{n}$  - Update affects:  $\sqrt{n}$  elements (one block) - Extra space:  $\sqrt{n}$  block minimums

**When to Use:** - Need both queries and updates - Array size up to  $10^6$  elements - Can't afford  $O(n \log n)$  space - Updates are as common as queries

## Algorithm 5: LCA-based RMQ (The Tree Transformation!)

### The Amazing Connection

RMQ can be transformed into finding the Lowest Common Ancestor (LCA) in a tree! We build a special tree called a Cartesian Tree.

### Building the Cartesian Tree Step-by-Step

**Array:** [5, 2, 4, 7, 1, 3]

#### Rules for Cartesian Tree:

1. In-order traversal gives the original array
2. Parent is always smaller than children (min-heap property)

**Construction Process: Step 1:** Add 5

5

**Step 2:** Add 2 (smaller than 5, becomes new root)

```
  2
   \
    5
```

**Step 3:** Add 4 (larger than 2, smaller than 5)

```
  2
   \
    4
     \
      5
```

**Step 4:** Add 7 (larger than all)


```
  2
   \
    4
     \
      5
       \
        7
```

**Step 5:** Add 1 (smallest, becomes new root)


```
  1
 /  \
2    3
 \
  4
   \
    5
     \
```

Wait, that's not right! Let me rebuild properly:

**Step 5:** Add 1 (smallest so far, becomes new root)



**Step 6:** Add 3 (larger than 1, goes to right)




### Final Cartesian Tree

Array: [5, 2, 4, 7, 1, 3]

Index: 0 1 2 3 4 5

Tree:



### How RMQ becomes LCA

**Key Insight:** The minimum element in range [L, R] is the LCA of nodes at positions L and R!

**Example: Query(1, 3) - Find min between indices 1 and 3**

1. Find nodes at indices 1 and 3 in the tree
  - Index 1 → Node with value 2
  - Index 3 → Node with value 7



2. Find their LCA:
  - Path from 2 to root:  $2 \rightarrow 1$
  - Path from 7 to root:  $7 \rightarrow 4 \rightarrow 2 \rightarrow 1$
  - Common ancestor: 2
3. Answer: 2

## LCA using Binary Lifting

We store ancestors at powers of 2 distances:

### Ancestor Table

Node	Parent	$2^1$ -ancestor	$2^2$ -ancestor
1	null	null	null
2	1	null	null
5	2	1	null
4	2	1	null
7	4	2	null
3	1	null	null

## Time Complexity Deep Dive

**Building the Cartesian Tree:  $O(n)$**  The amazing linear-time construction using a stack:

```
stack<int> st;
for (int i = 0; i < n; i++) {           // n iterations
    int last_popped = -1;
    while (!st.empty() && arr[st.top()] > arr[i]) { // amortized O(1)
        last_popped = st.top();
        st.pop();
    }
    // Set parent-child relationships - all O(1)
    if (!st.empty()) right_child[st.top()] = i;
    if (last_popped != -1) left_child[i] = last_popped;
    st.push(i);
}
```

**Why is the while loop  $O(1)$  amortized?** - Each element is pushed exactly once:  $n$  pushes total  
 - Each element is popped at most once:  $n$  pops total - Total operations across all iterations:  $2n = O(n)$  - Amortized per iteration:  $O(1)$

## LCA Preprocessing: $O(n \log n)$ for Binary Lifting Building the ancestor table:

For each node ( $n$  nodes):

For each power of 2 up to  $\log n$ :

$ancestor[node][j] = ancestor[ancestor[node][j-1]][j-1]$

Total:  $n \times \log n$  entries  $\times O(1)$  per entry =  $O(n \log n)$

Detailed breakdown:

Level 0: Store immediate parents  $\rightarrow n \times 1 = n$  operations  
 Level 1: Store 2-hop ancestors  $\rightarrow n \times 1 = n$  operations  
 Level 2: Store 4-hop ancestors  $\rightarrow n \times 1 = n$  operations  
 ...  
 Level  $\log n$ : Store  $2^{(\log n)}$  ancestors  $\rightarrow n \times 1 = n$  operations

Total:  $n \times (\log n + 1) = O(n \log n)$

**Query Time:  $O(\log n)$  for Binary Lifting LCA** Finding LCA of nodes  $u$  and  $v$ :

1. Bring  $u$  and  $v$  to same depth:
  - Calculate depths  $\rightarrow O(1)$  if preprocessed
  - Jump up using binary lifting  $\rightarrow O(\log n)$  jumps maximum
2. Binary search for LCA:
  - for  $j$  from  $\log n$  down to 0:  $\rightarrow \log n$  iterations
  - if  $\text{ancestor}[u][j] \neq \text{ancestor}[v][j]$ :
    - $u = \text{ancestor}[u][j]$
    - $v = \text{ancestor}[v][j]$   $\rightarrow O(1)$  per iteration

LCA = parent[ $u$ ]  $\rightarrow O(1)$

Total:  $O(\log n) + O(\log n) = O(\log n)$

**Example with depth difference = 13:**

13 in binary = 1101  
 Jump by:  $2^3 + 2^2 + 2^0 = 8 + 4 + 1 = 13$   
 Number of jumps: 3 (number of 1s in binary)  
 Maximum jumps:  $\log n$

**Alternative:  $O(1)$  Query with Euler Tour + RMQ** Preprocessing steps: 1. Build Cartesian Tree:  $O(n)$  2. Euler Tour of tree:  $O(n)$  - visit each edge twice 3. Build RMQ on depths:  $O(n)$  for  $\pm 1$  RMQ Total:  $O(n)$

**Query:** 1. Find first occurrence of  $u$  and  $v$  in tour:  $O(1)$  with preprocessing 2. RMQ on depth array between them:  $O(1)$  with  $\pm 1$  RMQ Total:  $O(1)$

This gives us  $O(n)$  preprocessing and  $O(1)$  query!

**Space Complexity Analysis Binary Lifting approach:**

- Cartesian Tree:  $n$  nodes  $\times$  3 pointers =  $3n$  pointers
- Ancestor table:  $n$  nodes  $\times$   $\log n$  levels =  $n \log n$  entries
- Depth array:  $n$  integers

Total:  $O(n \log n)$

**Euler Tour approach:**

- Cartesian Tree:  $3n$  pointers
- Euler tour:  $2n - 1$  entries

- Depth array:  $2n - 1$  entries
- First occurrence:  $n$  entries
- $\pm 1$  RMQ structure:  $O(n)$

Total:  $O(n)$

**Why Transform RMQ to LCA? Theoretical importance:** - Shows RMQ LCA (equivalent problems) - Any LCA solution gives RMQ solution - Any RMQ solution gives LCA solution - Unifies two seemingly different problems

**Practical benefits:** - Reuse existing LCA code - Some LCA variants are easier to solve - Opens door to other tree algorithms

**Real-world Performance** For  $n = 100,000$ : - Cartesian tree construction:  $\sim 10$  ms - LCA preprocessing:  $\sim 15$  ms - Total preprocessing:  $\sim 25$  ms - Query time:  $\sim 0.2$  microseconds - Memory:  $\sim 8$  MB

**Comparison with direct approaches:**

Algorithm	Preprocessing	Query	Theoretical Interest
-----	-----	-----	-----
Sparse Table	12 ms	0.05 s	Low
LCA-based	25 ms	0.20 s	Very High

The LCA approach is slightly slower but demonstrates beautiful theoretical connections!

**When to Use:** - You already have LCA code - Working with tree-related problems - Need to understand RMQ-LCA equivalence - Academic/competitive programming context

## Comparison: Which Algorithm to Choose?

### Quick Decision Guide

Need fast updates?

YES → Use Naive (no preprocessing) or Block Decomposition

NO → Continue...

Need  $O(1)$  queries?

YES → Use Sparse Table (best) or DP (if  $n < 1000$ )

NO → Continue...

Array size  $> 10000$ ?

YES → Use Sparse Table or Block Decomposition

NO → Use DP (simplest for small arrays)

### Performance Summary Table

Algorithm	Build Time	Query Time	Space	Best For
Naive	$O(1)$	$O(n)$	$O(n)$	Rare queries, frequent updates
DP	$O(n^2)$	$O(1)$	$O(n^2)$	Small arrays, many queries
Sparse Table	$O(n \log n)$	$O(1)$	$O(n \log n)$	Static arrays, fastest queries
Block Decomp	$O(n)$	$O(\sqrt{n})$	$O(\sqrt{n})$	Balanced operations
LCA-based	$O(n \log n)$	$O(\log n)$	$O(n \log n)$	Theoretical interest

## Interactive Examples

### Let's Trace Through a Complete Example

**Array:** [3, 1, 4, 1, 5, 9, 2, 6]

**Query(2, 6) = Find minimum in range [2, 6] Naive Approach:**

Check index 2: value = 4, min = 4

Check index 3: value = 1, min = 1

Check index 4: value = 5, min = 1

Check index 5: value = 9, min = 1

Check index 6: value = 2, min = 1

Answer: 1

### DP Approach:

Look up  $dp[2][6] = 1$  (pre-calculated)

Answer: 1

### Sparse Table:

Range length = 5

Largest power of 2  $\leq 5 = 4$

Query ranges: [2,5] and [3,6]

$st[2][2] = 1$ ,  $st[3][2] = 1$

Answer:  $\min(1, 1) = 1$

### Block Decomposition (block\_size = 3):

Blocks: [3,1,4] [1,5,9] [2,6,-]

Query spans partial block 0, complete block 1, partial block 2

Partial 0:  $\min(4) = 4$

Complete 1:  $block\_min = 1$

Partial 2:  $\min(2) = 2$

Answer:  $\min(4, 1, 2) = 1$

## Practice Problems

### Problem 1: Build Your Own DP Table

Array: [4, 2, 3, 1]

Fill in the DP table:

	j→	0	1	2	3
i↓		-----			
0		?	?	?	?
1		-	?	?	?
2		-	-	?	?
3		-	-	-	?

Solution

	j→	0	1	2	3
i↓		-----			
0		4	2	2	1
1		-	2	2	1
2		-	-	3	1
3		-	-	-	1

### Problem 2: Sparse Table Query

Given sparse table for array [6, 2, 5, 1, 7, 3]:

st[0][0]=6, st[0][1]=2, st[0][2]=1  
st[1][0]=2, st[1][1]=2, st[1][2]=1  
st[2][0]=5, st[2][1]=1, st[2][2]=1  
st[3][0]=1, st[3][1]=1  
st[4][0]=7, st[4][1]=3  
st[5][0]=3

What is Query(1, 4)?

Solution

Range length = 4, use  $k = 2$  ( $2^2 = 4$ ) Query ranges: [1,4] covered by st[1][2] = 1 Answer: 1

## Tips and Tricks

### 1. Sparse Table Power-of-2 Trick

Use bit operations for fast power-of-2 calculations:

```
int k = __builtin_clz(1) - __builtin_clz(range_length);  
// or  
int k = floor(log2(range_length));
```

### 2. Block Size Selection

For Block Decomposition, optimal block size is usually: -  $\sqrt{n}$  for balanced operations - Smaller blocks for faster queries - Larger blocks for faster updates

### 3. DP Memory Optimization

Only need to store the upper triangle of the DP table since  $dp[i][j]$  only makes sense when  $i \leq j$ .

### 4. Cartesian Tree Stack Trick

Build Cartesian Tree in  $O(n)$  using a stack to maintain the right spine of the tree.

## Deep Dive: Understanding Complexity Growth


### How Complexities Compare As N Grows

Let's see how each complexity grows with input size:

n	log n	$\sqrt{n}$	n	n log n	$n^2$
10	3	3	10	33	100
100	7	10	100	664	10,000
1,000	10	32	1,000	10,000	1,000,000
10,000	13	100	10,000	133,000	100,000,000
100,000	17	316	100,000	1,700,000	10,000,000,000
1,000,000	20	1,000	1,000,000	20,000,000	1,000,000,000,000

### Time Complexity Visualization

Operations vs Input Size (log scale)



### When Each Algorithm Wins

#### Scenario 1: Few Queries (< 100)

Total Time = Preprocessing + (Number of Queries  $\times$  Query Time)

For 50 queries on  $n = 100,000$ :

- Naive:  $0 + 50 \times 1000 \text{ s} = 50,000 \text{ s}$  (Winner!)
- Sparse Table:  $20,000 \text{ s} + 50 \times 0.05 \text{ s} = 20,002 \text{ s}$
- Block:  $100 \text{ s} + 50 \times 1 \text{ s} = 150 \text{ s}$

#### Scenario 2: Many Queries (> 1,000,000)

For 1,000,000 queries on  $n = 100,000$ :

- Naive:  $0 + 1\text{M} \times 1000 \text{ s} = 1,000,000,000 \text{ s}$  (16 minutes!)
- Sparse Table:  $20,000 \text{ s} + 1\text{M} \times 0.05 \text{ s} = 70,000 \text{ s}$  (Winner!)
- Block:  $100 \text{ s} + 1\text{M} \times 1 \text{ s} = 1,000,100 \text{ s}$

#### Scenario 3: Queries with Updates

For 10,000 queries + 1,000 updates on  $n = 100,000$ :

- Naive:  $10,000 \times 1000 \text{ s} + 1,000 \times 0 = 10,000,000 \text{ s}$



- Block:  $10,000 \times 1 \text{ s} + 1,000 \times 1 \text{ s} = 11,000 \text{ s}$  (Winner!)
- Sparse Table: Must rebuild after each update = Terrible!

## Memory vs Speed Trade-offs

Algorithm	Memory	Query Speed	Can Update?
Naive	4MB	Slow	Instant
DP	40GB	Fastest	Rebuild all
Sparse Table	6.4MB	Fastest	Rebuild all
Block	4MB	Fast	Fast
LCA	8MB	Fast	Rebuild tree

## Big-O Doesn't Tell the Whole Story!

**Hidden Constants Matter** Two  $O(n)$  algorithms can differ by 100x in practice:

```
// Algorithm A: O(n) with small constant
for (int i = 0; i < n; i++)
    sum += arr[i]; // 1 operation per iteration

// Algorithm B: O(n) with large constant
for (int i = 0; i < n; i++) {
    result = complex_hash(arr[i]); // 50 operations
    result = expensive_check(result); // 30 operations
    sum += result; // 80 operations total
}
```

Both are  $O(n)$ , but B is 80x slower!

**Cache Performance** Modern CPUs have cache hierarchies: - L1 Cache: 0.5 ns access time - L2 Cache: 7 ns access time - Main Memory: 100 ns access time

Sequential access (cache-friendly):

`arr[0], arr[1], arr[2], ...` → All from cache!

Random access (cache-unfriendly):

`arr[1000], arr[0], arr[5000], ...` → Cache misses!

This is why Sparse Table (sequential access) often beats theoretically faster algorithms with random access patterns.

## Amortized Analysis: When Average Case Matters

### Example: Dynamic Array Resizing

Push operations: 1, 1, 1, (resize+copy:4), 1, 1, 1, 1, (resize+copy:8), ...

Individual operations:  $O(1)$  usually,  $O(n)$  sometimes

Amortized (average):  $O(1)$  per operation!

**In RMQ Context** Block Decomposition updates: - Best case: New min is smaller  $\rightarrow O(1)$  - Worst case: Recompute block  $\rightarrow O(\sqrt{n})$  - Amortized: Often  $O(1)$  in practice

### The Complexity Hierarchy

$O(1)$     $O(\log n)$     $O(\sqrt{n})$     $O(n)$     $O(n \log n)$     $O(n^2)$     $O(2^n)$

Constant < Logarithmic < Sublinear < Linear < Linearithmic < Quadratic < Exponential

**Rule of thumb for max input sizes:** -  $O(1)$ ,  $O(\log n)$ : Any size (limited by memory) -  $O(\sqrt{n})$ : Up to  $10^{14}$  -  $O(n)$ : Up to  $10^8$  -  $O(n \log n)$ : Up to  $10^6$  -  $O(n^2)$ : Up to  $10^4$  -  $O(n^3)$ : Up to 500 -  $O(2^n)$ : Up to 20

### Conclusion

Each RMQ algorithm represents a different trade-off between preprocessing time, query time, space usage, and update capability. Understanding not just the big-O notation but also:

- Hidden constants in the implementation
- Cache performance characteristics
- Amortized vs worst-case behavior
- Memory access patterns
- Practical input size limits

...helps you choose the right algorithm for your specific use case.

Remember: - **Naive** = No prep, just scan (best for rare queries) - **DP** = Pre-calculate everything (impractical for large arrays) - **Sparse Table** = Smart power-of-2 ranges (best for static arrays) - **Blocks** = Divide and conquer (best with updates) - **LCA** = Transform to tree problem (theoretical elegance)

The “best” algorithm depends entirely on your specific requirements!

Happy querying!